




Figure 7. Reliability diagrams on each channel's native scale, with no post-hoc fitting: the integer and expected-value scores are divided by 100, the log-odds channel goes through a logistic. The dotted line is perfect calibration. This measures whether a deployment could read these scores as probabilities as-is.